

# BANGIRANA DOUGLAS /00573

1 LOGISTIC REGRESSION

In [1]:

```
1 #Required Libraries
2 from sklearn.datasets import make_classification
3 from sklearn.model_selection import GridSearchCV
4 from sklearn.linear_model import LogisticRegression
5 from sklearn.model_selection import train_test_split
6 import pandas as pd
7 import numpy as np
```

In [2]:

```
1 df=pd.read_csv("C:\\Users\\Atwongire Vianney\\Desktop\\AI_PRAC\\loan_data.csv")
2 df
```

Out[2]:

	Gender	Married	Dependents	Education	ApplicantIncome	CoapplicantIncome	LoanAmount
0	Male	Yes	1	Graduate	4583	1508.0	128
1	Male	Yes	0	Graduate	3000	0.0	66
2	Male	Yes	0	Not Graduate	2583	2358.0	120
3	Male	No	0	Graduate	6000	0.0	141
4	Male	Yes	0	Not Graduate	2333	1516.0	95
...	...	...	...	...	...	...	...
376	Male	Yes	3+	Graduate	5703	0.0	128
377	Male	Yes	0	Graduate	3232	1950.0	108
378	Female	No	0	Graduate	2900	0.0	71
379	Male	Yes	3+	Graduate	4106	0.0	40
380	Female	No	0	Graduate	4583	0.0	133

381 rows × 9 columns



In [3]: 1 df.head()

Out[3]:

	Gender	Married	Dependents	Education	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Status
0	Male	Yes	1	Graduate	4583	1508.0	128	1
1	Male	Yes	0	Graduate	3000	0.0	66	1
2	Male	Yes	0	Not Graduate	2583	2358.0	120	1
3	Male	No	0	Graduate	6000	0.0	141	1
4	Male	Yes	0	Not Graduate	2333	1516.0	95	1

In [4]: 1 df.tail()

Out[4]:

	Gender	Married	Dependents	Education	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Status
376	Male	Yes	3+	Graduate	5703	0.0	128	1
377	Male	Yes	0	Graduate	3232	1950.0	108	1
378	Female	No	0	Graduate	2900	0.0	71	1
379	Male	Yes	3+	Graduate	4106	0.0	40	1
380	Female	No	0	Graduate	4583	0.0	133	1

In [5]: 1 #checking for missing values  
2 print(df.isnull().sum())

```
Gender          5
Married         0
Dependents      8
Education       0
ApplicantIncome 0
CoapplicantIncome 0
LoanAmount      0
Loan_Amount_Term 11
Loan_Status     0
dtype: int64
```

```
In [6]: 1 #Handling missing values appropriately
2 filled_df=df.fillna(0)
3 print(filled_df)
4
5 #Dropping rows with missing values
6 dropped_df=df.dropna()
7 print(dropped_df)
8
9 #saving the cleaned dataset to a new csv file
10 filled_df.to_csv('cleaned_file.csv', index=False)
11 dropped_df.to_csv('cleaned_file_drop.ped.csv', index=False)
```

	Gender	Married	Dependents	Education	ApplicantIncome	\
0	Male	Yes	1	Graduate	4583	
1	Male	Yes	0	Graduate	3000	
2	Male	Yes	0	Not Graduate	2583	
3	Male	No	0	Graduate	6000	
4	Male	Yes	0	Not Graduate	2333	
..	...	...	...	...	...	
376	Male	Yes	3+	Graduate	5703	
377	Male	Yes	0	Graduate	3232	
378	Female	No	0	Graduate	2900	
379	Male	Yes	3+	Graduate	4106	
380	Female	No	0	Graduate	4583	

	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Loan_Status
0	1508.0	128	360.0	N
1	0.0	66	360.0	Y
2	2358.0	120	360.0	Y
3	0.0	141	360.0	Y
4	1516.0	95	360.0	Y
..	...	...	...	...
376	0.0	128	360.0	Y
377	1950.0	108	360.0	Y
378	0.0	71	360.0	Y
379	0.0	40	180.0	Y
380	0.0	133	360.0	N

[381 rows x 9 columns]

	Gender	Married	Dependents	Education	ApplicantIncome	\
0	Male	Yes	1	Graduate	4583	
1	Male	Yes	0	Graduate	3000	
2	Male	Yes	0	Not Graduate	2583	
3	Male	No	0	Graduate	6000	
4	Male	Yes	0	Not Graduate	2333	
..	...	...	...	...	...	
376	Male	Yes	3+	Graduate	5703	
377	Male	Yes	0	Graduate	3232	
378	Female	No	0	Graduate	2900	
379	Male	Yes	3+	Graduate	4106	
380	Female	No	0	Graduate	4583	

	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Loan_Status
0	1508.0	128	360.0	N
1	0.0	66	360.0	Y
2	2358.0	120	360.0	Y
3	0.0	141	360.0	Y
4	1516.0	95	360.0	Y
..	...	...	...	...
376	0.0	128	360.0	Y
377	1950.0	108	360.0	Y
378	0.0	71	360.0	Y
379	0.0	40	180.0	Y
380	0.0	133	360.0	N

[358 rows x 9 columns]

```
In [7]: 1 df.isnull()
```

```
Out[7]:
```

	Gender	Married	Dependents	Education	ApplicantIncome	CoapplicantIncome	LoanAmount
0	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False
...	...	...	...	...	...	...	...
376	False	False	False	False	False	False	False
377	False	False	False	False	False	False	False
378	False	False	False	False	False	False	False
379	False	False	False	False	False	False	False
380	False	False	False	False	False	False	False

381 rows × 9 columns




```
In [8]: 1 df_cleaned = df.dropna()  
2 df_cleaned
```

```
Out[8]:
```

	Gender	Married	Dependents	Education	ApplicantIncome	CoapplicantIncome	LoanAmount
0	Male	Yes	1	Graduate	4583	1508.0	128
1	Male	Yes	0	Graduate	3000	0.0	66
2	Male	Yes	0	Not Graduate	2583	2358.0	120
3	Male	No	0	Graduate	6000	0.0	141
4	Male	Yes	0	Not Graduate	2333	1516.0	95
...	...	...	...	...	...	...	...
376	Male	Yes	3+	Graduate	5703	0.0	128
377	Male	Yes	0	Graduate	3232	1950.0	108
378	Female	No	0	Graduate	2900	0.0	71
379	Male	Yes	3+	Graduate	4106	0.0	40
380	Female	No	0	Graduate	4583	0.0	133

358 rows × 9 columns



```
In [9]: 1 filled_df.to_csv('cleaned_file.csv', index=False)
```

```
In [10]: 1 x=df_cleaned['LoanAmount'].array.reshape(-1,1)
        2 x
```

```
Out[10]: <PandasArray>
```

```
[
  [128],
  [66],
  [120],
  [141],
  [95],
  [70],
  [109],
  [114],
  [17],
  [125],
  [100],
  [76],
  [133],
  [104],
  [116],
  [122],
  [110],
  ...]
```

```
In [11]: 1 y=df_cleaned['Loan_Amount_Term']
        2 y
```

```
Out[11]: 0      360.0
        1      360.0
        2      360.0
        3      360.0
        4      360.0
        ...
       376      360.0
       377      360.0
       378      360.0
       379      180.0
       380      360.0
        Name: Loan_Amount_Term, Length: 358, dtype: float64
```

```
In [12]: 1 #splitting the data
        2 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_st.
```

```
In [13]: 1 x_train.shape
```

```
Out[13]: (286, 1)
```







```
In [23]: 1 import pandas as pd
2 from sklearn.linear_model import LinearRegression
3 from sklearn.model_selection import train_test_split
```

```
In [24]: 1 df.head()
```

Out[24]:

	Gender	Married	Dependents	Education	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Status
0	Male	Yes	1	Graduate	4583	1508.0	128	1
1	Male	Yes	0	Graduate	3000	0.0	66	1
2	Male	Yes	0	Not Graduate	2583	2358.0	120	1
3	Male	No	0	Graduate	6000	0.0	141	1
4	Male	Yes	0	Not Graduate	2333	1516.0	95	1

```
In [25]: 1 df.tail()
```

Out[25]:

	Gender	Married	Dependents	Education	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Status
376	Male	Yes	3+	Graduate	5703	0.0	128	1
377	Male	Yes	0	Graduate	3232	1950.0	108	1
378	Female	No	0	Graduate	2900	0.0	71	1
379	Male	Yes	3+	Graduate	4106	0.0	40	1
380	Female	No	0	Graduate	4583	0.0	133	1

```
In [26]: 1 df.isnull().sum()
```

```
Out[26]: Gender          5
Married          0
Dependents       8
Education        0
ApplicantIncome  0
CoapplicantIncome 0
LoanAmount       0
Loan_Amount_Term 11
Loan_Status      0
dtype: int64
```



```
In [29]: 1 y=df_cleaned['Loan_Amount_Term']
        2 y
```

```
Out[29]: 0      360.0
        1      360.0
        2      360.0
        3      360.0
        4      360.0
        ...
        376    360.0
        377    360.0
        378    360.0
        379    180.0
        380    360.0
        Name: Loan_Amount_Term, Length: 358, dtype: float64
```

```
In [30]: 1 #splitting the data
        2 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_st
```

```
In [31]: 1 model=LinearRegression()
        2 model.fit(x_train,y_train)
```

```
Out[31]: ▾ LinearRegression
          LinearRegression()
```

```
In [32]: 1 score=model.score(x_train,y_train)
        2 score
```

```
Out[32]: 0.013044972022888901
```

```
In [33]: 1 model.coef_
```

```
Out[33]: array([0.28778211])
```

```
In [34]: 1 model.intercept_
```

```
Out[34]: 308.4210137519316
```

## FINE TUNING

```
In [35]: 1 #Defining parameters
        2 from sklearn.model_selection import GridSearchCV
        3 from sklearn.metrics import mean_squared_error
```

```
In [36]: 1 param_grid={
2         'copy_X':[True,False],
3         'fit_intercept':[True,False],
4         'n_jobs':[True,False],
5         'positive':[True,False]
6
7     }
```

```
In [37]: 1 #performing grid with cross_validation
2 grid_search=GridSearchCV(model,param_grid,cv=5,scoring='neg_mean_squared_e
3 grid_search.fit(x_train,y_train)
```

```
Out[37]:  ▸ GridSearchCV
          ▸ estimator: LinearRegression
              ▸ LinearRegression
```

```
In [38]: 1 #Best model
2 best_model=grid_search.best_estimator_
3 best_model
```

```
Out[38]:  ▾ LinearRegression
          LinearRegression(n_jobs=True, positive=True)
```

```
In [39]: 1 best_score=model.score(x_test,y_test)
2 best_score
```

```
Out[39]: -0.009395552322493339
```

```
In [40]: 1 y_pred=best_model.predict(x_test)
2 y_pred
```

```
Out[40]: array([345.83268771, 341.8037382 , 337.19922449, 348.13494456,
346.12046981, 339.21369924, 343.53043085, 331.73136445,
340.07704556, 337.19922449, 345.5449056 , 333.45805709,
327.70241494, 342.37930242, 325.40015808, 346.12046981,
321.65899069, 341.5159561 , 342.95486663, 339.50148134,
348.71050878, 347.27159824, 323.67346544, 342.95486663,
343.81821295, 342.95486663, 334.32140341, 345.25712349,
321.9467728 , 347.84716246, 340.07704556, 344.39377717,
325.9757223 , 348.13494456, 334.32140341, 349.86163721,
349.86163721, 338.92591713, 340.65260977, 342.95486663,
336.04809606, 342.95486663, 351.01276564, 325.11237598,
347.27159824, 342.95486663, 338.35035291, 328.56576127,
347.55938035, 345.25712349, 346.40825192, 340.94039188,
345.83268771, 315.61556644, 332.59471077, 345.83268771,
348.99829089, 342.95486663, 329.71688969, 348.13494456,
346.12046981, 336.33587816, 337.19922449, 328.56576127,
321.08342647, 337.19922449, 350.43720142, 327.70241494,
340.94039188, 330.0046718 , 340.07704556, 345.83268771])
```

```
In [41]: 1 mse=mean_squared_error(y_test,y_pred)
          2 mse
```

```
Out[41]: 2688.917596325753
```

```
In [ ]: 1
```